

Pick any publicly-traded company that trades on the Nasdaq or the NYSE.

What company did you select, and what is its ticker symbol?

Apple Inc., which trades on the Nasdaq under the ticker symbol "AAPL".

B

Load the data

```
In [ ]: #C
import pandas as pd

apple=pd.read_csv("AAPL.csv",index_col="Date",parse_dates=True)

Out[ ]:
      Date      Open      High      Low      Close  Adj Close  Volume
2022-04-29  162.11998  165.16193  162.04307  163.28005  161.89501  9604400
2022-04-28  162.20000  162.33996  156.72001  156.90003  155.85581  9662320
2022-04-27  159.92004  158.78993  155.38025  155.97007  155.82728  8803200
2022-04-26  159.20000  158.52004  156.82993  163.39999  162.86494  13021800
2022-04-25  161.83904  164.16997  167.25000  167.68994  166.70079  13174700
2022-04-22  166.72007  168.22995  162.29004  167.82007  167.06866  12059260
2022-04-21  158.14904  160.71007  156.52007  159.47999  158.51930  8696560
2022-04-20  159.68988  166.47995  159.29995  166.02004  165.02070  10625600
2022-04-19  163.88000  164.08002  164.86997  156.77004  155.82690  13062300
2022-04-18  156.00995  159.44002  154.17993  157.27999  156.36268  118124600
```

```
In [ ]: #C
apple.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 248 entries, 2022-04-25 to 2023-04-21
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   Date        248 non-null    datetime64[ns]
 1   Open        248 non-null    float64
 2   High        248 non-null    float64
 3   Low         248 non-null    float64
 4   Close       248 non-null    float64
 5   Adj Close   248 non-null    float64
 6   Volume      248 non-null    int64  
dtypes: datetime64[ns]: 1, float64: 5
memory usage: 13.7 KB
```

```
In [ ]: #D
# check the data type of the index column
apple.index.dtype

Out[ ]:
dtype('datetime64[ns]')
```

D

The DataFrame provided is indexed by time values, as evidenced by the "Date" column and confirmed by checking the data type of the index column.

E(A)

Now, view the max and min value of your index attribute.

```
In [ ]: #E(A)
apple.index.max()

Out[ ]:
Timestamp('2023-04-21 09:00:00')

In [ ]: #E(A)
apple.index.min()

Out[ ]:
Timestamp('2022-04-25 09:00:00')
```

E(B)

Now, view the argmax and argmin values of your index attribute.

```
In [ ]: #E(B)
apple.index.argmax()

Out[ ]:
249


In [ ]: #E(B)
apple.index.argmin()

Out[ ]:
0

What do the results of max, min, argmax, and argmin represent?
max and min: These are the maximum and minimum values in the index, respectively. In a time-based index, these values represent the latest and earliest datetimes in the DataFrame.
argmax and argmin: These are the indices (e.g., row numbers) that correspond to the maximum and minimum values in the index, respectively. In a time-based index, these values represent the row numbers that correspond to the latest and earliest datetimes in the DataFrame.
So in the context of the DataFrame provided, if the max and min values of the index are '2023-04-20' and '2022-04-25', respectively, then the DataFrame covers a period of time from April 25, 2022 to April 21, 2023. And if the argmax and argmin values of the index are 249 and 0, respectively, then the latest datetime is in row 249 of the DataFrame, and the earliest datetime is in row 0 of the DataFrame.
```

```
In [ ]: #E(F)
apple.plot()

Out[ ]:
<Axes: xlabel='Date'>
```




F(a)

The plot shows multiple lines, it may be challenging to interpret because it can be difficult to distinguish between the different lines and to determine which line corresponds to which variable. Additionally, the plot may be difficult to read if the x-axis labels are not formatted in a clear and consistent manner, or if there are too many data points or too much information displayed at once. To make the plot easier to understand, it can be helpful to add labels to the axes and to the individual lines, to use different colors or line styles for different lines, and to adjust the scale and formatting of the axes to better display the data. Additionally, it may be helpful to plot subsets of the data as a line or to aggregate the data in a meaningful way before plotting it. As soon as its going to plot for all the axes and its really important to choose the ones that makes sense and helpful to the model as said above.

```
In [ ]: #E(F)
apple['Close'].plot()

Out[ ]:
<Axes: xlabel='Date'>
```

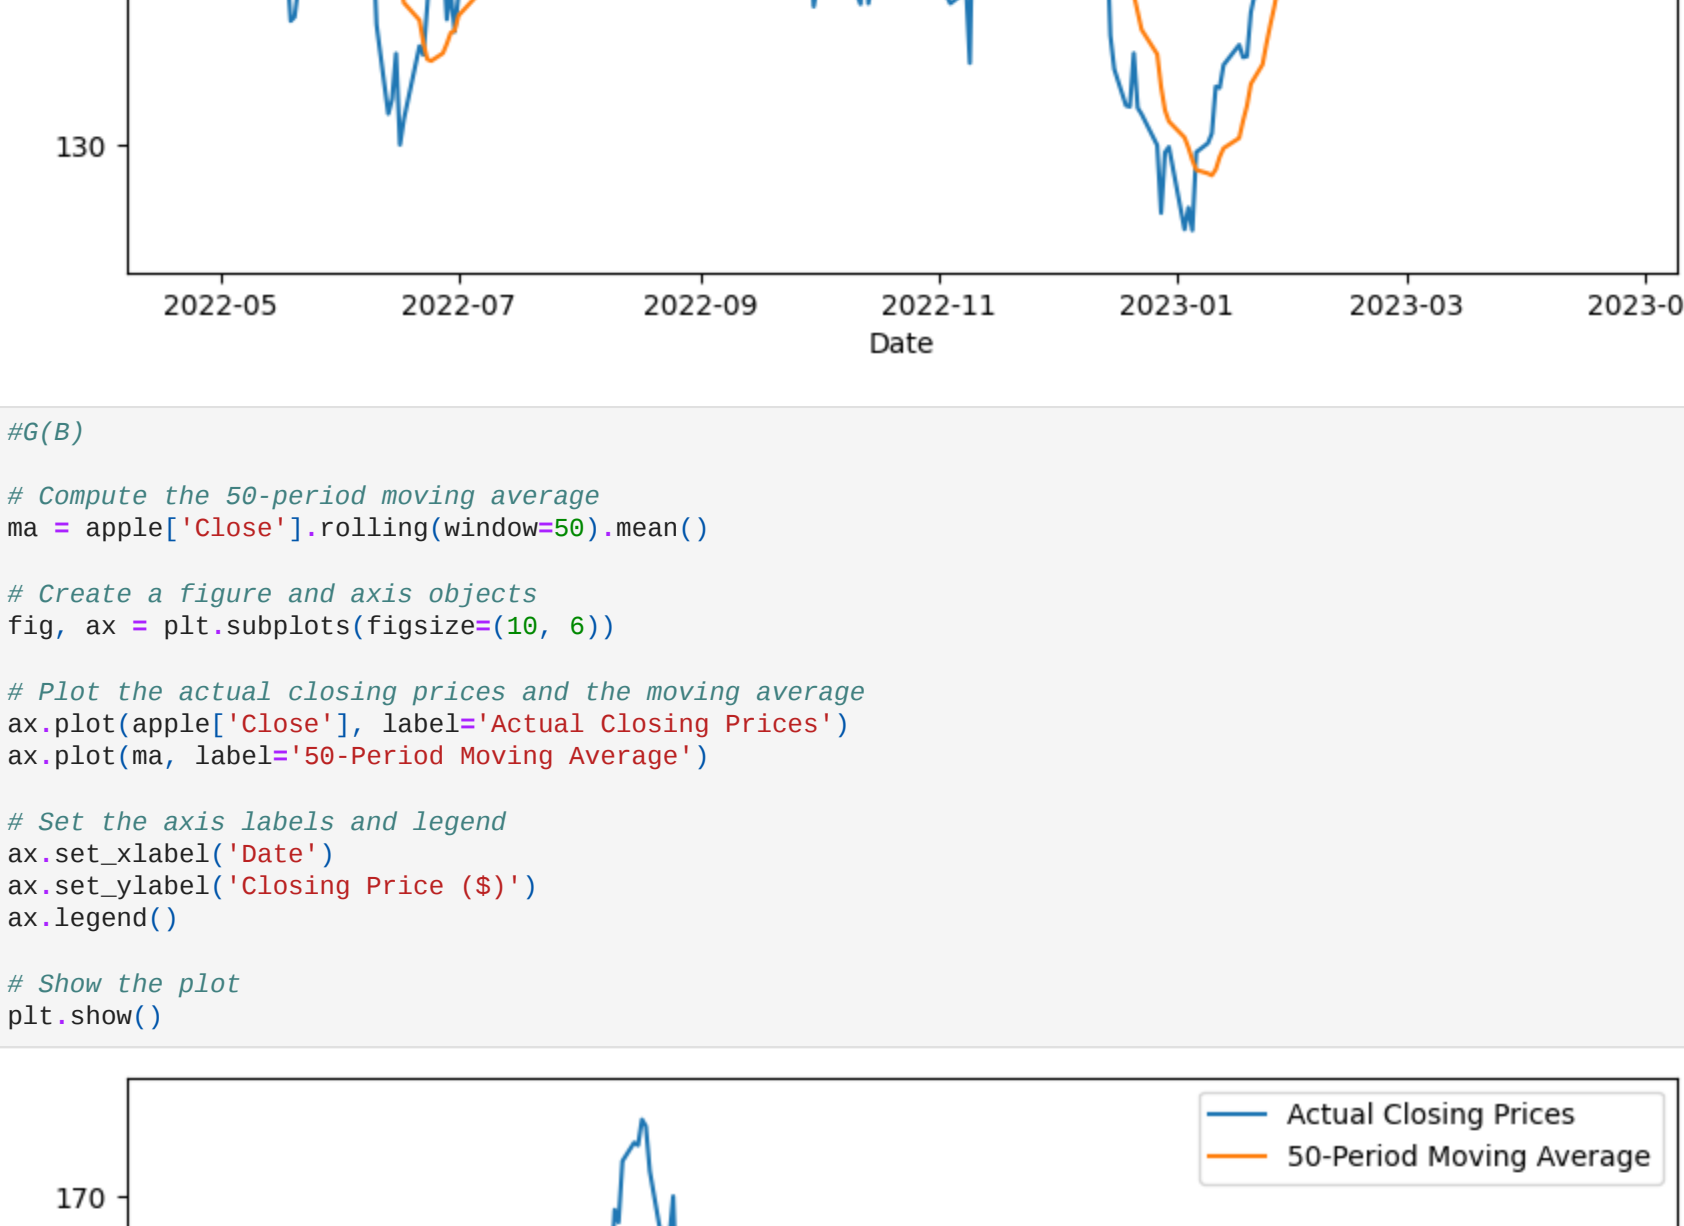


F(B)

Plotting the 'Close' variable only makes the graph more easily interpretable because it simplifies the visual representation to only show the closing prices for the stock. This allows for a clearer understanding of the overall trend of the stock price over time, without the potential confusion of having multiple lines for different variables.

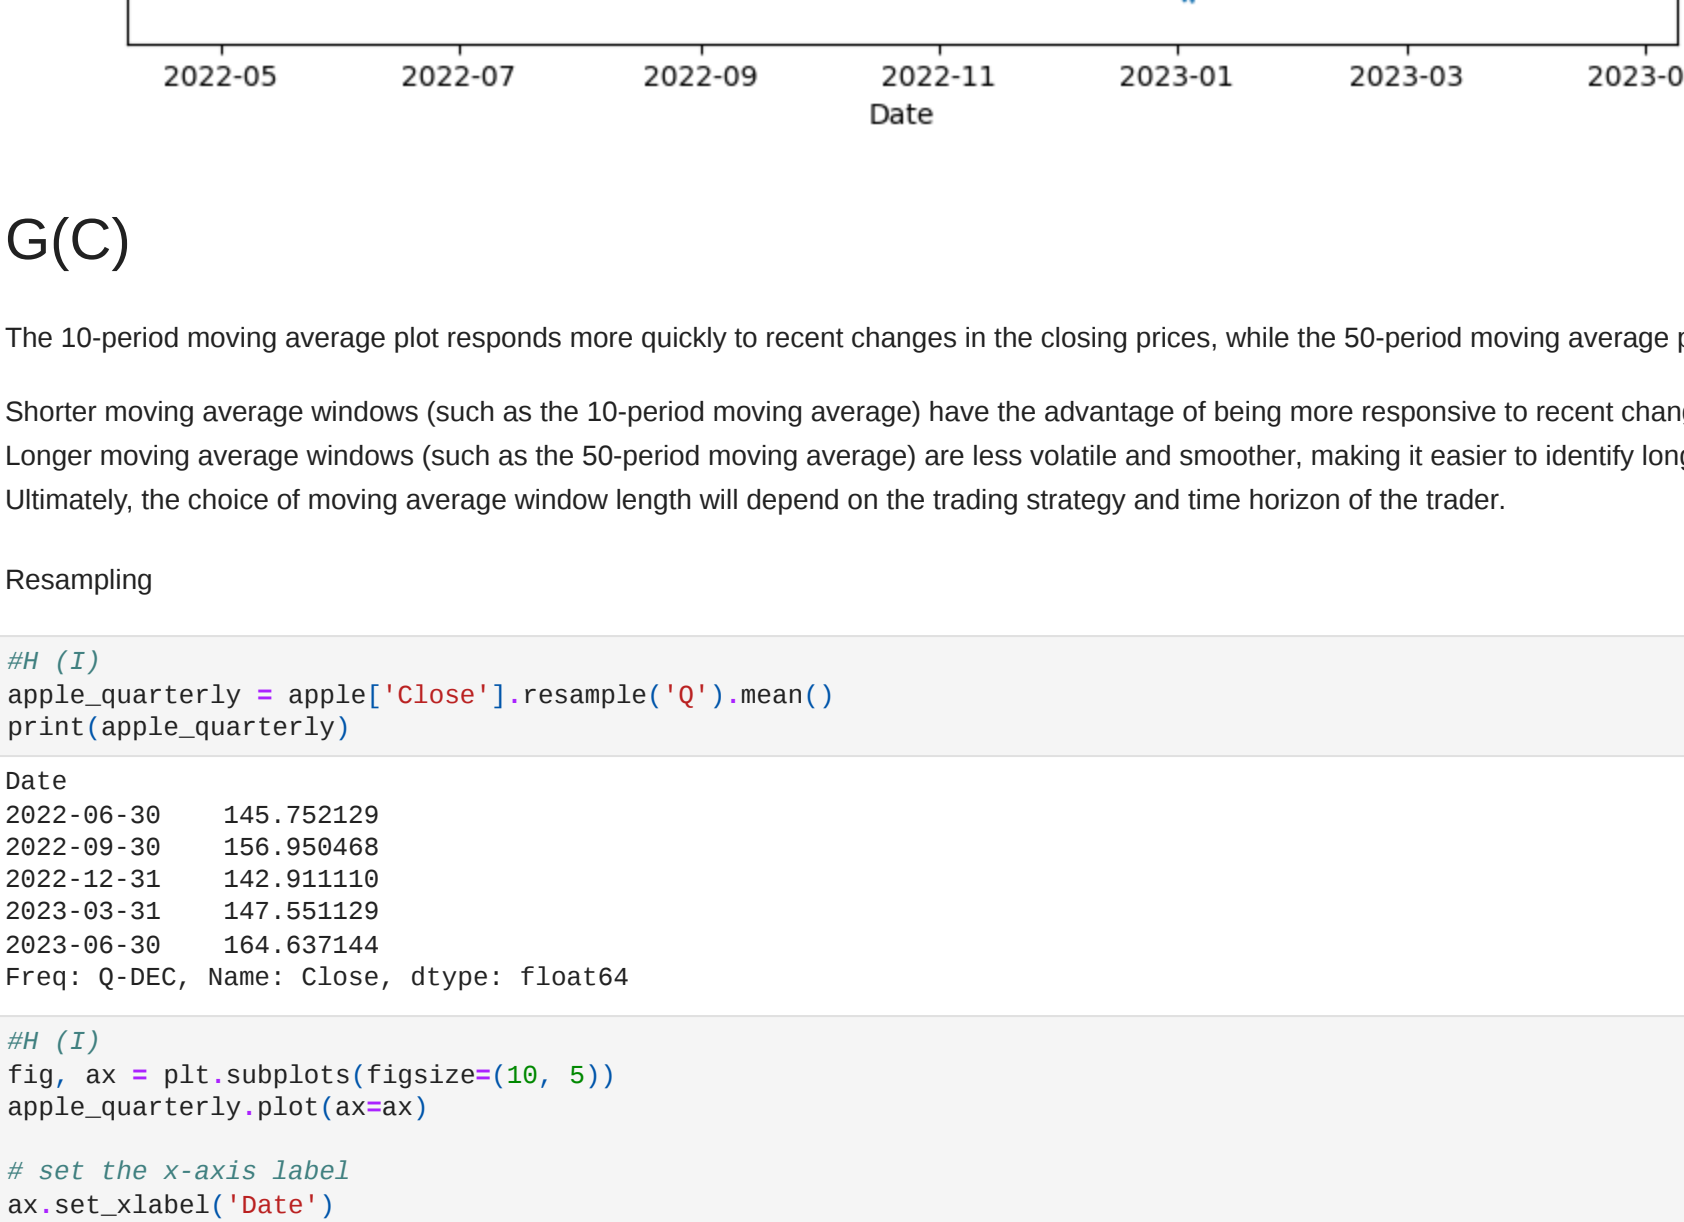
```
In [ ]: #E(F)
apple['2022-04-25':'2022-04-28']['Close'].plot()

Out[ ]:
<Axes: xlabel='Date'>
```



```
In [ ]: #E(F)
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
apple['2022-04-25':'2022-04-28']['Close'].plot(color='red', linestyle='--')
plt.title('Apple Stock Prices in April 2022')
plt.xlabel('Date')
plt.ylabel('Closing Price ($)')
plt.show()
```



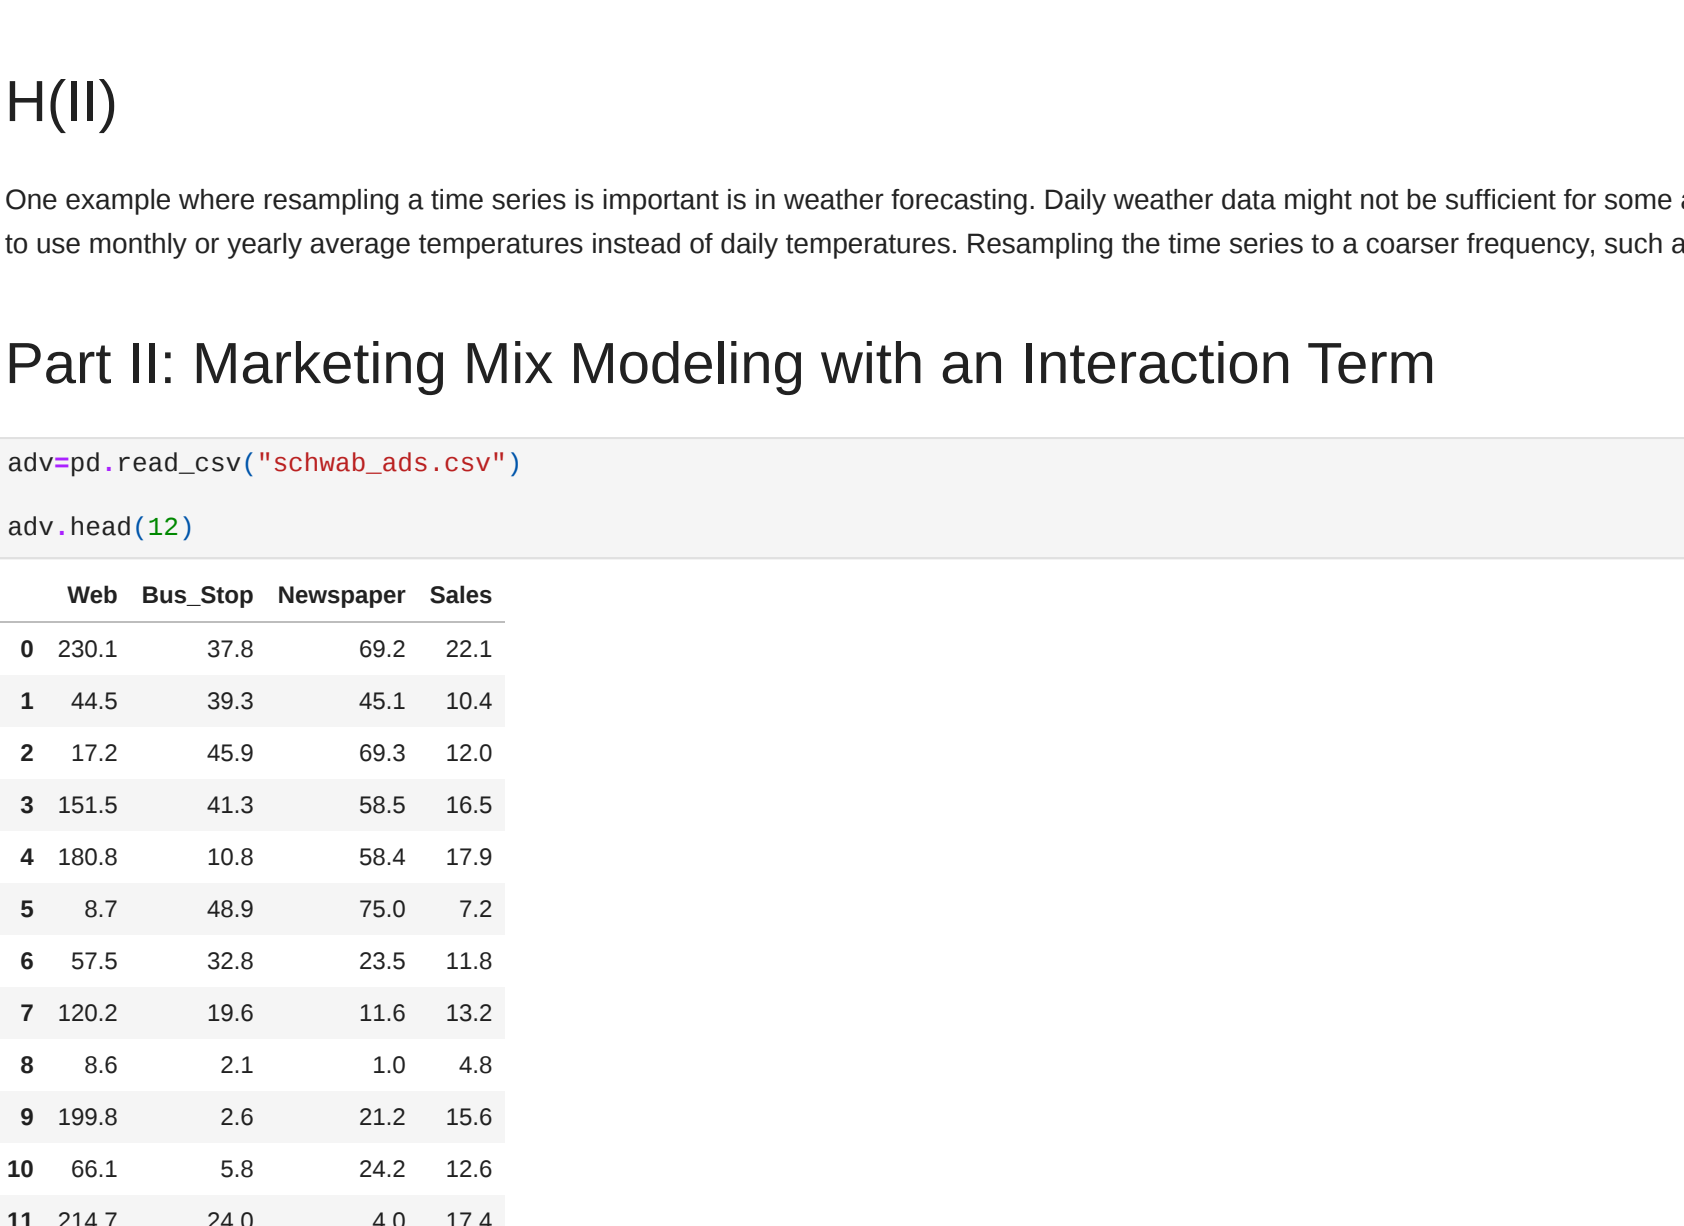
```
In [ ]: #E(F)
# Compute the 10-period moving average
ma = apple['Close'].rolling(window=10).mean()

# Create a figure and axis objects
fig, ax = plt.subplots(figsize=(10, 6))

# Plot the actual closing prices and the moving average
ax.plot(apple['Close'], label='Actual Closing Prices')
ax.plot(ma, label='10-Period Moving Average')

# Set the axis labels and legend
ax.set_xlabel('Date')
ax.set_ylabel('Closing Price ($)')
ax.legend()

# Show the plot
plt.show()
```



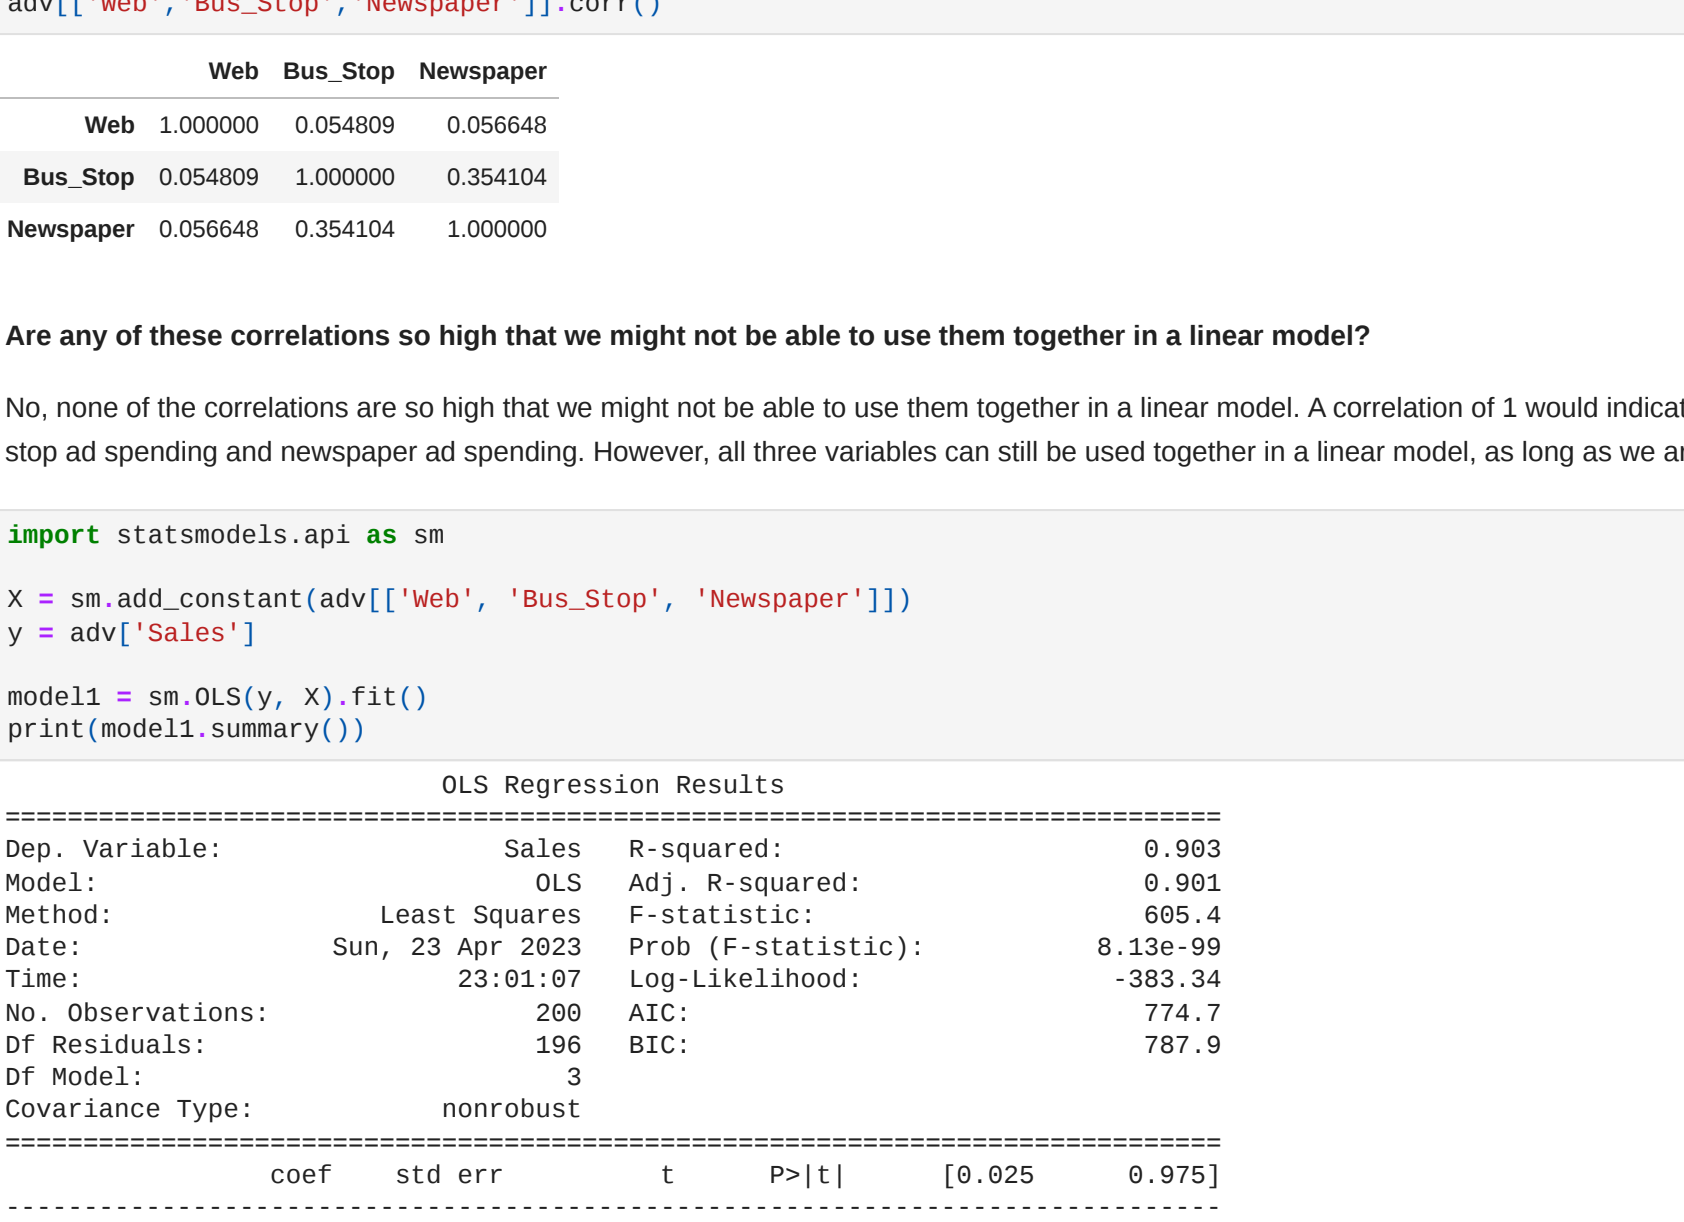
```
In [ ]: #E(F)
# Compute the 50-period moving average
ma = apple['Close'].rolling(window=50).mean()

# Create a figure and axis objects
fig, ax = plt.subplots(figsize=(10, 6))

# Plot the actual closing prices and the moving average
ax.plot(apple['Close'], label='Actual Closing Prices')
ax.plot(ma, label='50-Period Moving Average')

# Set the axis labels and legend
ax.set_xlabel('Date')
ax.set_ylabel('Closing Price ($)')
ax.legend()

# Show the plot
plt.show()
```



G(C)

The 10-period moving average plot responds more quickly to recent changes in the closing prices, while the 50-period moving average plot responds more slowly and is smoother.

Shorter moving average windows (such as the 10-period moving average) have the advantage of being more responsive to recent changes in prices, allowing traders to identify trends earlier. However, they are also more volatile and susceptible to false signals, which can lead to losses. Longer moving average windows (such as the 50-period moving average) are less volatile and smoother, making it easier to identify long-term trends. However, they are less responsive to recent changes in prices, which can result in missed opportunities to enter or exit trades. Ultimately, the choice of moving average window length will depend on the trading strategy and time horizon of the trader.

```
In [ ]: #E(F)
apple.quarterly = apple['Close'].resample('Q').mean()
print(apple.quarterly)

date
2022-06-30    145.762129
2022-09-30    156.950468
2022-12-31    145.931118
2023-03-31    147.551129
2023-06-30    164.837444
Freq: Q-DEC, Name: Close, dtype: float64
```

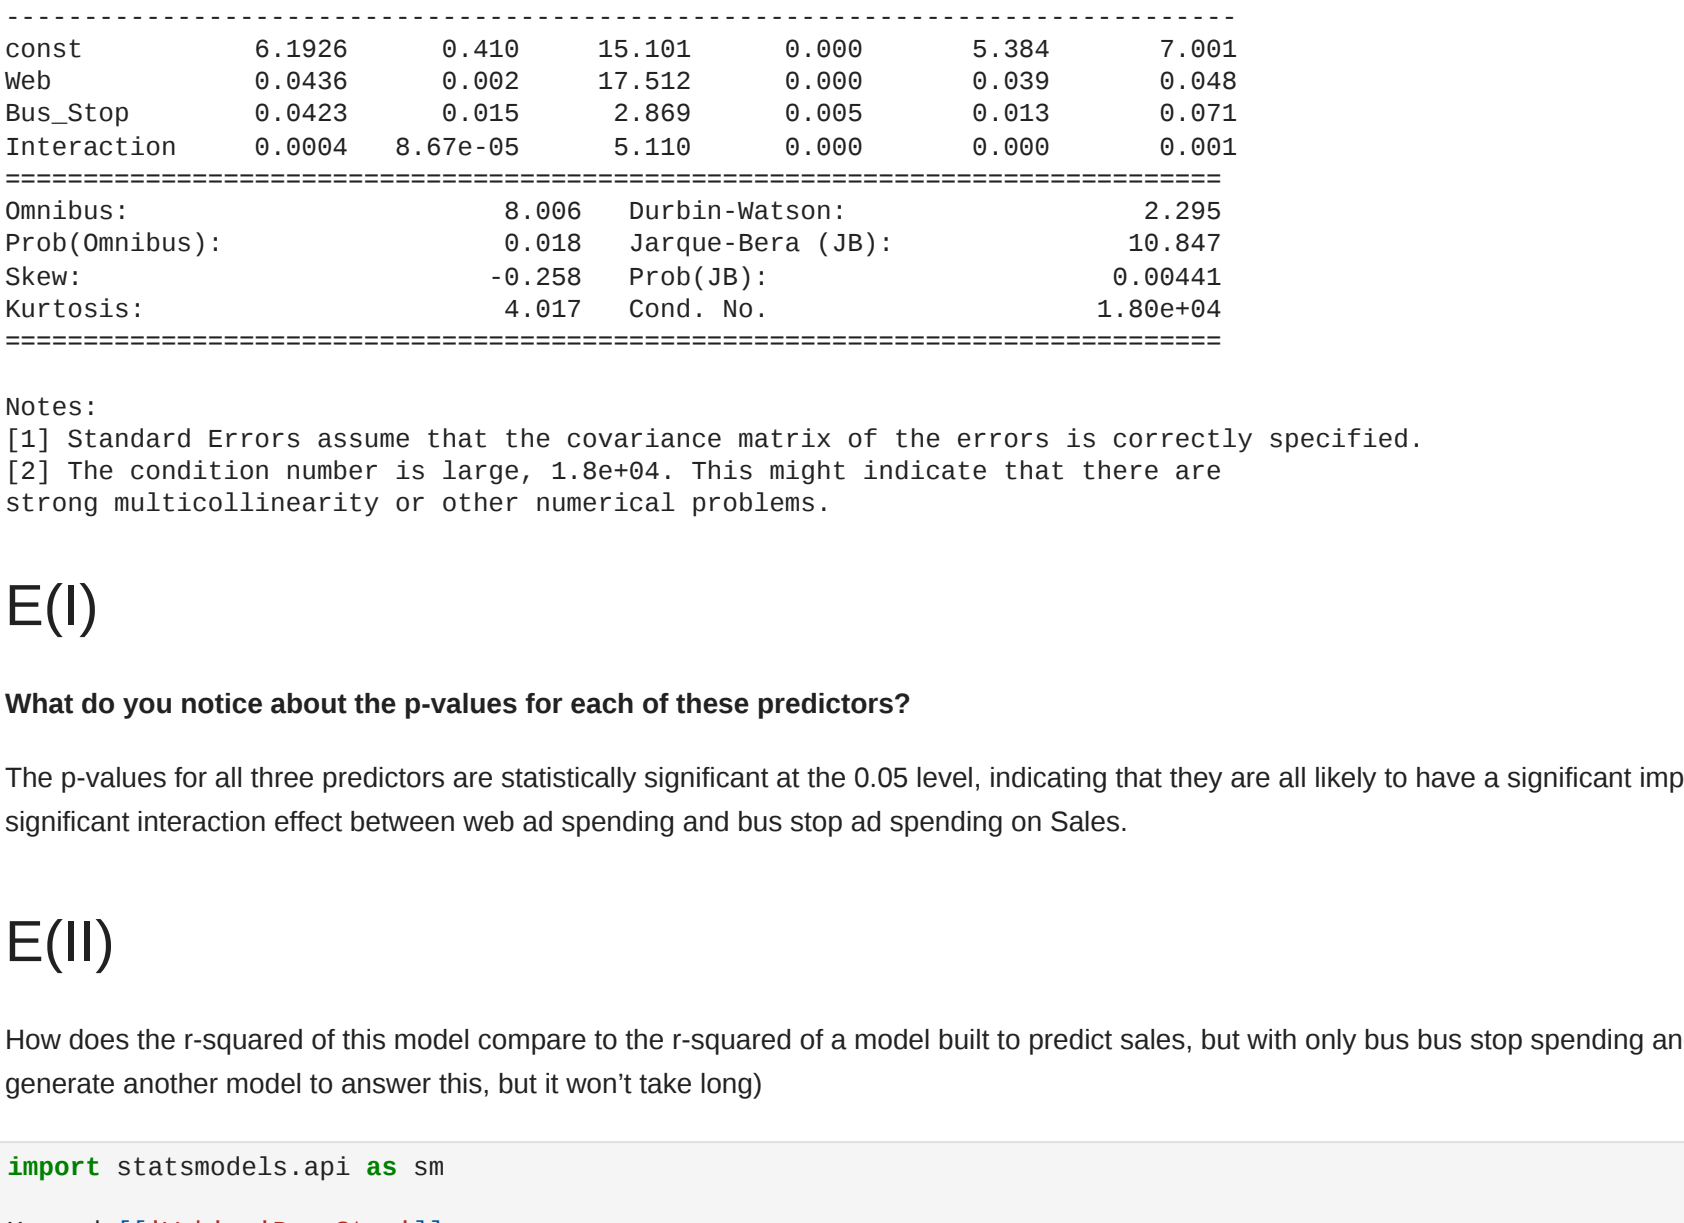
```
In [ ]: #E(F)
fig, ax = plt.subplots(figsize=(10, 6))
apple.quarterly.plot(ax=ax)

# Set the x-axis label
ax.set_xlabel('Date')

# Set the y-axis label
ax.set_ylabel('Closing Price (USD)')

# Add the plot title
ax.set_title('Apple Quarterly Closing Prices')

# Show the plot
plt.show()
```



H(I)

One example where resampling a time series is important is in weather forecasting. Daily weather data might not be sufficient for some analyses or applications. For example, if you wanted to assess the trend in average temperature over a year or a decade, it would be more appropriate to use monthly or yearly average temperatures instead of daily temperatures. Resampling the time series to a coarser frequency, such as monthly or yearly, can help smooth out short-term fluctuations and reveal longer-term patterns in the data.

Part II: Marketing Mix Modeling with an Interaction Term

```
In [ ]: adv=pd.read_csv("schwab_ads.csv")

adv.head(12)

   Web  Bus_Stop  Newspaper  Sales
0  29.1    37.8    69.2    2.1
1  44.5    36.3    45.1    16.4
2  17.2    45.9    69.3    12.0
3  15.5    41.3    58.5    16.5
4  18.8    10.8    58.4    7.0
5   8.7    48.9    75.0    7.2
6  57.8    32.8    21.2    11.8
7  120.2   18.6    11.6    13.2
8   8.6     2.1    1.0     4.8
9  199.8   25.6    21.2    15.6
10  66.1    5.8    24.2    12.6
11  234.7   34.0    4.0    17.4
```

```
In [ ]: # create a new variable "Total_Spending" that shows the total spending
adv['Total_Spending'] = adv['web'] + adv['Bus_Stop'] + adv['Newspaper']

adv.head(10)

   Web  Bus_Stop  Newspaper  Sales  Total_Spending
0  29.1    37.8    69.2    2.1          136.1
1  44.5    36.3    45.1    16.4          126.0
2  17.2    45.9    69.3    12.0          132.4
3  15.5    41.3    58.5    16.5          115.3
4  18.8    10.8    58.4    7.0          87.0
5   8.7    48.9    75.0    7.2          131.8
6  57.8    32.8    21.2    11.8          111.8
7  120.2   18.6    11.6    13.2          151.4
8   8.6     2.1    1.0     4.8           11.7
9  199.8   25.6    21.2    15.6          246.6
```

```
In [28]: adv['Total_Spending'].corr(adv['Sales'])

Out[28]:
0.92491706249931

The high correlation value of 0.92 suggests a strong positive linear relationship between total spending and sales. However, correlation does not imply causation. There may be other factors that influence sales, such as market demand, competition, product quality, pricing, etc. Also, the relationship may not be linear or may have a time lag. Therefore, we cannot conclude from this correlation alone that more ad spending leads to more sales. Further analysis, such as regression modeling, controlled experiments, and domain expertise, may be needed to establish causal relationships and make actionable recommendations.
```

```
In [29]: #C
adv[['Web','Bus_Stop','Newspaper']].corr()

Out[29]:
          Web  Bus_Stop  Newspaper
Web  1.000000  0.054009  0.056648
Bus_Stop  0.054009  1.000000  0.354004
Newspaper  0.056648  0.354004  1.000000
```

Are any of these correlations so high that we might not be able to use them together in a linear model?

No, none of the correlations are so high that we might not be able to use them together in a linear model. A correlation of 1 would indicate a perfect linear relationship, but the highest correlation value here is only 0.354104, which suggests a relatively weak linear relationship between bus stop ad spending and newspaper ad spending. However, all three variables can still be used together in a linear model, as long as we are aware of their respective correlations and interpret the results accordingly.

```
In [30]: import statsmodels.api as sm

X = adv[['Web','Bus_Stop']]
y = adv['Sales']

model1 = sm.OLS(y, X).fit()
print(model1.summary())

OLS Regression Results
=====
Dep. Variable:      Sales      R-squared:      0.003
Model:              OLS      Adj. R-squared:    0.002
Method:             Least Squares      F-statistic:  665.4
Date:               Sun, 23 Apr 2023      Prob (F-statistic):  2.39e-180
Time:               23:01:12      Log-Likelihood: -383.34
No. Observations:    280          AIC:          746.7
Df Residuals:        198          BIC:          787.9
Df Model:             2
Covariance Type:     nonconstant
=====
               coef      std err      t      P>|t|      [0.025   0.975]
-----
const          5.6375     0.481    11.720     0.000     4.680     6.595
Web             0.0002     0.000     0.000     1.000     -0.000     0.000
Bus_Stop        0.1872     0.000    19.726     0.000     0.185     0.189
Interaction      0.0004     0.000     0.152     0.879     -0.001     0.001
=====
Omnibus:          2.183      Durbin-Watson:   2.262
Prob(Omnibus):    0.034      Jarque-Bera (JB):   1.848
Skew:             -0.147      P>Skew(B):    0.8643
Kurtosis:         4.617      Cond. No.     1.20e+04
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.
```

```
In [40]: import numpy as np

web = 220
bus_stop = 30

# coefficients from model
coef = np.array([4.637, 0.004, 0.1872])

# calculate predicted sales
sales_pred = np.dot(coef, np.array([1, web, bus_stop]))

print('Predicted Sales: ', sales_pred)

Predicted Sales: 19.8149
```

E(IV)

The interaction effect between bus stop ad and bus stop ad spending suggests that the combination of these two advertising methods has a greater impact on sales than the sum of their individual effects. In other words, the effect of web ad spending on sales depends on the level of bus stop ad spending and vice versa. The predicted sales for a marketer using 220 units of web ad spending and 30 units of bus stop ad spending is 19.8149.

```
In [40]: #E
import statsmodels.api as sm

X = adv[['Web','Bus_Stop','Total_Spending']]
X['Web_Total'] = X['Web'] * X['Total_Spending']
X['Bus_Stop_Total'] = X['Bus_Stop'] * X['Total_Spending']
y = adv['Sales']

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())

OLS Regression Results
=====
Dep. Variable:      Sales      R-squared:      0.033
Model:              OLS      Adj. R-squared:    0.030
Method:             Least Squares      F-statistic:  1.61e+111
Date:               Mon, 24 Apr 2023      Prob (F-statistic):  1.00e-111
Time:               06:05:23      Log-Likelihood: -346.52
No. Observations:    280          AIC:          705.6
Df Residuals:        198          BIC:          724.8
Df Model:             5
Covariance Type:     nonconstant
=====
               coef      std err      t      P>|t|      [0.025   0.975]
-----
const          5.6375     0.481    11.720     0.000     4.680     6.595
Web             0.0002     0.000     0.000     1.000     -0.000     0.000
Bus_Stop        0.1872     0.000    19.726     0.000     0.185     0.189
Web_Total       -0.0002     0.000    -0.491     0.627     -0.001     0.000
Bus_Stop_Total  -0.0002     0.000    -0.957     0.340     -0.001     0.000
=====
Omnibus:          2.183      Durbin-Watson:   2.262
Prob(Omnibus):    0.034      Jarque-Bera (JB):   1.848
Skew:             -0.147      P>Skew(B):    0.8643
Kurtosis:         4.617      Cond. No.     1.20e+04
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.
```

```
In [40]: #E
import statsmodels.api as sm

X = adv[['Web','Bus_Stop','Total_Spending']]
X['Web_Total'] = X['Web'] * X['Total_Spending']
X['Bus_Stop_Total'] = X['Bus_Stop'] * X['Total_Spending']
y = adv['Sales']

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
print(model.summary())

OLS Regression Results
=====
Dep. Variable:      Sales      R-squared:      0.033
Model:              OLS      Adj. R-squared:    0.030
Method:             Least Squares      F-statistic:  1.61e+111
Date:               Mon, 24 Apr 2023      Prob (F-statistic):  1.00e-111
Time:               06:05:23      Log-Likelihood: -346.52
No. Observations:    280          AIC:          705.6
Df Residuals:        198          BIC:          724.8
Df Model:             5
Covariance Type:     nonconstant
=====
               coef      std err      t      P>|t|      [0.025   0.975]
-----
const          5.6375     0.481    11.720     0.000     4.680     6.595
Web             0.0002     0.000     0.000     1.000     -0.000     0.000
Bus_Stop        0.1872     0.000    19.726     0.000     0.185     0.189
Web_Total       -0.0002     0.000    -0.491     0.627     -0.001     0.000
Bus_Stop_Total  -0.0002     0.000    -0.957     0.340     -0.001     0.000
=====
Omnibus:          2.183      Durbin-Watson:   2.262
Prob(Omnibus):    0.034      Jarque-Bera (JB):   1.848
Skew:             -0.147      P>Skew(B):    0.8643
Kurtosis:         4.617      Cond. No.     1.20e+04
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.
```

Part II: Wildcards: A Real-World Consulting Opportunity

One possible recommendation to us as an online platform for grading assignments, such as Blackboard or Turnitin. This would allow the lecturer to easily delegate grading to TAs or even peer-reviewers, while maintaining consistency in grading standards. Another option is to implement self-grading or peer grading assignments, where students are given rubrics to grade each other's work, with the lecturer serving as a final arbiter for any discrepancies. This can also provide students with valuable feedback and improve their understanding of course material. Additionally, implementing automated grading tools, such as Grammarly or Turnitin's Feedback Studio, can help reduce the time spent grading assignments, although it may not be appropriate for all types of assignments. Finally, the lecturer can consider reducing the number of assignments or adjusting the workload to make grading more manageable.

Another additional point would be - using a combination of TA grading and peer grading could also be considered, where students grade each other's assignments under the supervision of the TA. This can be a good way to reduce the workload on the lecturer while also encouraging student engagement and feedback.